



IOPMP Task Group Meeting

April 11, 2024

[Video link](#)

Minutes

- Released draft 6, change log:
- Clarification: MDSTALL.is_stalled
- Issue: what if the rule of MDCFG.t is breached?

Draft 6, Change log

- Change log:
 - Rename SID to RRID
 - Per-entry interrupt control
 - Per-entry bus error control
 - Multi-fault record
 - Refine, clarify, and typo-fixup

Clarification: MDSTALL.is_stalled

- What if writing a non-zero value to MDSTALL, but no RRID is selected?
 - **MDSTALL.is_stalled** should be “1”, because
 - It can avoid infinite polling of this flag because it is never asserted.
 - That is, only writing a zero to MDSTALL(H) (resume command) can de-assert “**is_stalled.**”

A reference flow to program an IOPMP

Step 1.1: write MDSTALL once *// exactly once*

Step 1.2: write RRIDSCP zero or more times

Step 1.3: poll until MDSTALL.is_stalled == 1 *// ensure all stalls take effect*

Step 2: update IOPMP's configuration

Step 3.1: write MDSTALL=0 *// resume all transactions*

Step 3.2: poll until MDSTALL.is_stalled == 0 *// ensure all resumes take effect*

Some steps may be skipped according to implementations or use cases.

Issue: what if **MDCFG.t** is NOT monotonically incremental?

- In full-model or isolation-model, what if **MDCFG[j].t** > **MDCFG[j+1].t** when writing **MDCFG[j]** or **MDCFG[j+1]**?
 - Policy:
 - 1) Implementation-dependent, no specific HW behavior is required
 - 2) List one or a few options to implement
 - Possible options:
 - 1) Ignore the write causing the case // some programming order is required.
 - 2) For $j < m \leq 62$, fix **MDCFG[m].t** = MIN { **MDCFG[k].t** | $0 \leq k \leq m$ }
 - 3) Modify **MDCFG[j+1].t** = **MDCFG[j].t**. // Channing will post the slides separately
 - 4) Mark MD(j+1) to MD(62) empty (have no entry) // possible imp on next page

