



IOPMP Task Group Meeting


November 7, 2024

[Video link](#)

Minutes

- The feedback of r0.9.2-RC1:
 - wrong k -value position/# of bits: **fixed**
 - MSI support or not? **Yes, taking 3 new fields/registers**
 - clarification: SRCMD table fmt=2: **the spec need a figure to help readers**
 - SRCMD table fmt=2 for overwriting permission within an MD: **explanation**
 - SRCMD table fmt=2 for MDSTALL: **simplified implementation**

Adjust HWCFG0.rsv



For MDCFG_FMT=:

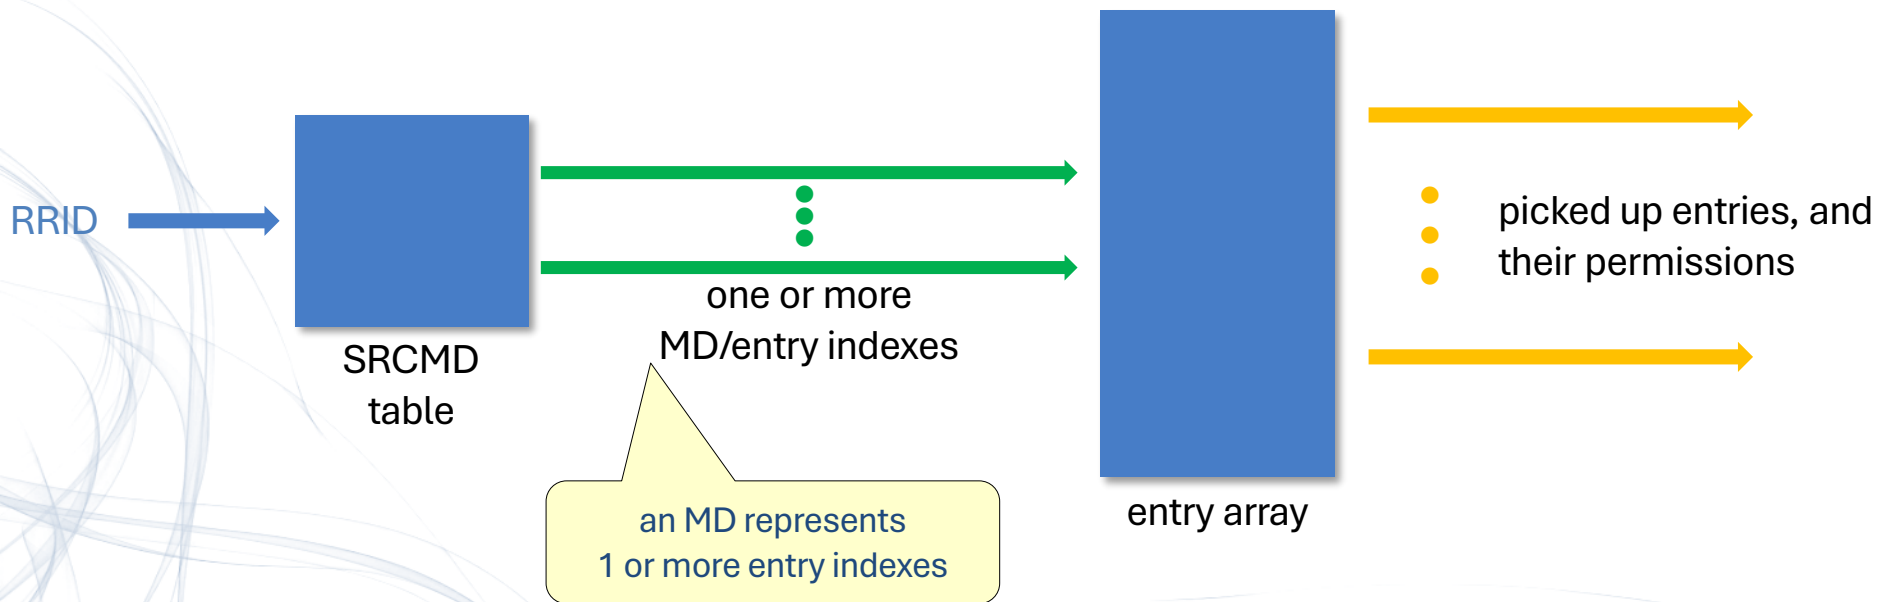
- 0: ZERO [0] for full model; //MD is configured by MDCFG
- 1: RO [k value - 1] for rapid- k ; //No MDCFG
- 2: RW/RO [k value - 1] for dynamic- k ; //No MDCFG

It can still support up to $128 \times 63 = 8,064$ entries for the rapid- k .
 For even larger requests, the full model can support up to 65,535 entries.

Shall the IOPMP spec support MSI?

- Current spec only defines a single interrupt signal, no MSI supported. To use MSI, one should use APLIC now.
- Shall we direct supporting MSI?
 - Registers needed:
 - A 11-bit or 32-bit MSI data to write
 - EIID at APLIC: 11 bits
 - IOMMU uses 32 bits
 - A 32-bit and 64-bit MSI address to write
 - A bit indicates which mode (single wire or MSI) is implemented.
 - A lock mechanism (maybe we can reuse ERR_CFG.I)

Entries pick up of IOPMP (SRCMD.fmt=0)



Format 0 vs Format 2

Format 0:
the table was index by RRID

RRID 0 0 0 0 ... 1 1 1

RRID 1 0 0 1 ... 0 0 1

⋮

RRID i 0 1 0 ... 0 0 0

MD/entry j
MD/entry 1
MD/entry 0

SRCMD
table

Format 2:
the table is index by entry

Entry 0 01 00 00 ... 00 00 11

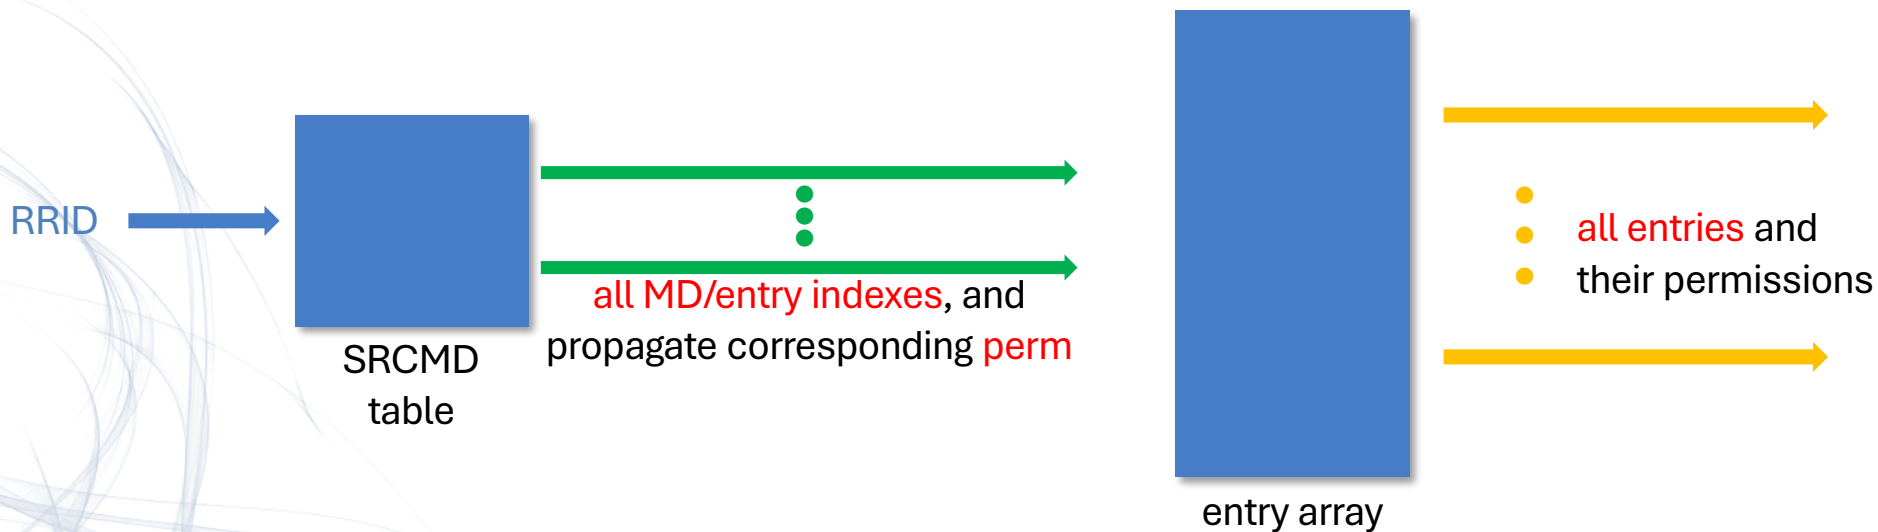
Entry 1 01 00 00 ... 00 11 01

Entry j 01 00 11 ... 00 00 00

write for RRID i
read for RRID i

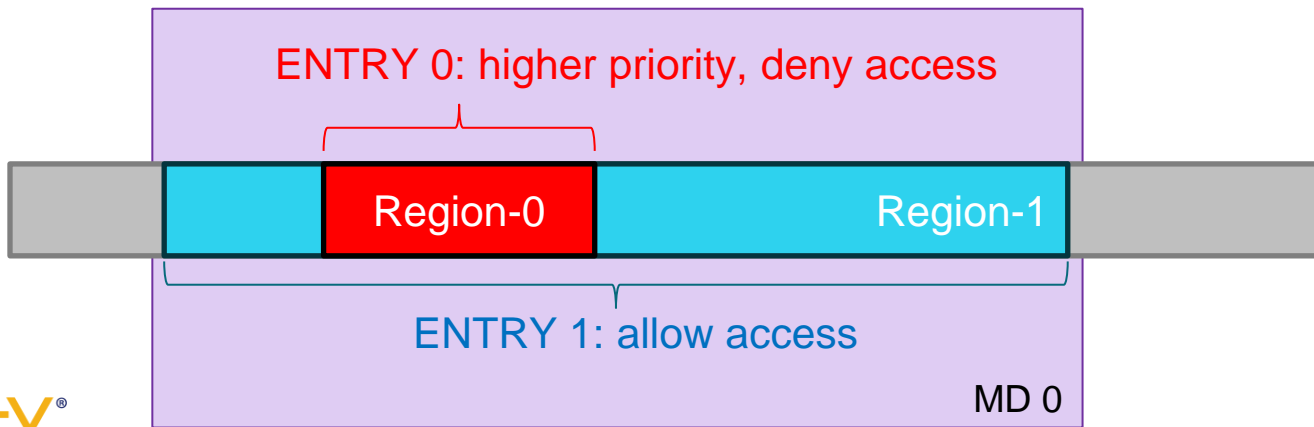
write for RRID 0
read for RRID 0

Entries pick up of IOPMP (SRCMD.fmt=2)



Permission exception within an MD

- Drill a permission exception by higher-priority rules within an MD, for `SRCMD.fmt ≠ 2`
- EX: `MD0 = (ENTRY 0, ENTRY 1)`
 - `ENTRY_CFG(0).r/w=0b00` for region 0 // narrow range, block access
 - `ENTRY_CFG(1).r/w=0b11` for region 1 // wider range, allow access



Priority rules under SRCMD.fmt = 2

- Drill a permission exception by higher-priority rules for SRCMD.fmt = 2:
- Problem: MD[0]= (ENTRY[0], ENTRY[1])
 - ENTRY_CFG(0).r/w=0b00 for region 0
 - ENTRY_CFG(1).r/w=0b11 for region 1
 - Add SRCMD(0).PERM[r_{rid}].r/w=0b00 //that is, let ENTRY_CFG decide

r/w permission for RRID r == PERM[r].r/w OR ENTRY_CFG.r/w

MDSTALL for SRCMD.fmt = 2

- MDSTALL.md stalls the related RRIDs by selecting MDs; however, every MD is related to all RRIDs. That is, all RRIDs will be stalled together as long as you select any MDs.
- Thus, in MDSTALL(H), only the bit 0 is effective : the field **exempt** (for write) or **is_stalled** (for read), no matter how many MDs are implemented.
- Proposal: in SRCMD.fmt=2, if one wants to implement MDSTALL(H), one can implement bit 0 in MDSTALL only no matter how many MDs are implemented.
 - Let MDSTALL \leftarrow 1 // stall all RRIDs
 - Let MDSTALL \leftarrow 0 // resume all RRIDs
- ✓ RRIDSCP: not need to change.